

READ: REINFORCEMENT-GUIDED CAPABILITY DISTILLATION FOR LARGE LANGUAGE MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

Capability distillation applies knowledge distillation to selected model capabilities, aiming to compress a large language model (LLM) into a smaller one while preserving the abilities needed for a downstream task. However, most existing methods treat capabilities as independent training targets and overlook how improving one capability can reshape the student’s broader capability profile, especially when multiple abilities jointly determine task success. We study capability distillation under a fixed token budget and identify two consistent patterns: distillation induces systematic, budget-dependent cross-capability transfer, and additional budget often brings limited task-relevant gains while sometimes degrading other useful abilities. Building on these insights, we propose *ReAD*, a *Reinforcement-guided cApability Distillation* framework that explicitly accounts for capability interdependence. ReAD first infers task-essential capabilities, then generates capability-targeted supervision on the fly, and finally uses an uncertainty-aware contextual bandit to adaptively allocate the distillation budget based on expected utility gains. Extensive experiments show that ReAD improves downstream utility under the same token budget while reducing harmful spillover and wasted distillation effort compared to strong baselines. Our code is available at <https://anonymous.4open.science/r/ReAD-62B7>.

1 INTRODUCTION

Knowledge distillation (KD) (Hinton et al., 2015) for large language models (LLMs) (Zhao et al., 2023) has become an essential research direction for enhancing the accessibility and efficiency of Machine-Learning-as-a-Service (MLaaS) (Cai et al., 2024; Xu et al., 2024; Yang et al., 2024): **a small LLM learns to imitate a large teacher’s outputs and can then be deployed under limited computational or financial resources** (Sanh et al., 2019; Jiao et al., 2020; Wang et al., 2020; Li et al., 2021; Liang et al., 2023). Recent studies show that capability distillation, a specialization of knowledge distillation that focuses supervision on a target capability (e.g., instruction-following, reasoning, mathematics, or coding), can substantially improve the student’s performance on downstream tasks that primarily depend on that capability, while achieving these gains at lower serving cost than the large teacher model (Magister et al., 2023; Shridhar et al., 2023; Xu et al., 2023a; Yue et al., 2024).

Despite rapid progress in capability distillation, most existing methods still treat capabilities as if they can be improved independently (Taori et al., 2023; Chiang et al., 2023; Zhang et al., 2024), even though empirical evidence suggests that optimizing one capability often triggers broad, unintended shifts in the model’s overall capability profile (Zhong et al., 2024; Fang et al., 2025; Cloud et al., 2025). This mismatch is especially consequential in the budgeted setting: under a fixed token budget, the goal is to maximize downstream task utility, which depends not only on the targeted capability but also on other capabilities the task may implicitly require. If cross-capability interactions are ignored, allocating more tokens to a single target can become inefficient in two ways: the extra improvement on the target can shrink as the budget grows, and the same updates may reduce performance on other tasks that require non-target capabilities. In this case, additional tokens are effectively spent on updates that provide little task-relevant benefit and may even increase deployment risk in MLaaS, which we refer to as *budget waste*. **To examine this gap, we run a controlled empirical study.** Concretely, we repeatedly distill the student under multiple token budgets, each time allocating the entire budget to one target capability drawn from a set of widely recognized core LLM capabilities: General Knowledge (General), Reasoning, Math, Code, Tool use, Long-Context Understanding

(LCU), Steerability, and Multilinguality. After each run, we evaluate the resulting student on the full benchmark suite in Table 3 and record the score change for every capability. **This yields a budget-dependent capability transfer matrix (diagonal: on-target gains; off-diagonal: redistribution across the other capabilities).** Across budgets, we observe two consistent patterns: (i) *capability-specific distillation induces systematic, budget-dependent transfer to other capabilities rather than isolated improvements*, and (ii) *increasing the budget for a single target produces smaller additional target gains while making the average harm to non-target capabilities more pronounced*, which together explain why naively scaling the budget for one capability can be inefficient.

Building on these observed insights, we propose *ReAD*, a *Reinforcement-guided cApability Distillation* framework for large language models. Overall, ReAD combines on-the-fly capability-targeted data generation with token-level knowledge distillation, and uses a lightweight contextual bandit to adaptively allocate a fixed budget across interdependent capabilities. Specifically, ReAD first infers a task requirement vector that identifies which capabilities are essential for improving the downstream utility and treats degradations on these dimensions as harmful spillover. It then allocates distillation effort to maximize a proxy reward that favors capability gains aligned with the task requirements while penalizing spillover and budget consumption. In each interval, ReAD samples capability-labeled prompts according to the current allocation, queries the teacher to obtain supervision, updates the student with a standard distillation loss, and uses the resulting capability-profile change to update an uncertainty-aware UCB allocation rule. Extensive experiments demonstrate that ReAD strengthens task-relevant capabilities and reduces wasted budget on low-utility capability updates under budget constraints.

To sum up, this paper makes the following key contributions:

- **Empirical study of cross-capability interactions.** We show that distilling toward a single capability consistently changes other capabilities, revealing systematic, **often overlooked** cross-capability interactions **that existing methods ignore**.
- **Budgeted capability distillation formulation.** We formulate capability distillation as allocating a fixed budget across multiple interacting capabilities, providing an objective **that improves utility while explicitly controlling side effects**.
- **ReAD: reinforcement-guided capability distillation.** ReAD infers task-relevant capabilities, generates capability-targeted distillation data with controllable style and difficulty, and uses an uncertainty-aware contextual bandit to allocate the budget across capabilities.
- **Theory analysis of the capability interdependence.** We explain cross-capability interactions and diminishing returns in capability distillation.

2 PRELIMINARY

2.1 NOTATION

In this paper, we study budgeted capability distillation from a large teacher model T to a smaller student model S , with total training-token budget B . Let $\mathcal{C} = \{c_1, \dots, c_{|\mathcal{C}|}\}$ denote a set of measurable capabilities, and let $s_k(M)$ be model M 's benchmark score on capability c_k . A distillation strategy π specifies how data is generated, how the teacher is queried, and how tokens are allocated across capabilities, assigning budgets $\{b_k\}$ with $\sum_k b_k \leq B$, and producing a distilled student S_π .

2.2 INTUITION AND FORMALIZATION

Distilling for one capability can inadvertently shift others, so a limited budget is easily spent inefficiently (Zhong et al., 2024; Fang et al., 2025; Cloud et al., 2025). To study these interactions, we evaluate eight core LLM capabilities with commonly adopted benchmarks (Appendix A, Table 3), **which motivate our unified capability distillation formulation**.

Definition 2.1 (Capability distillation). Given a teacher model T , an initial student model S_0 , and capability-specific data distributions $\{\mathcal{D}_c\}_{c \in \mathcal{C}}$, capability distillation allocates a limited budget B across capabilities to improve task-relevant performance under shared representations by optimizing a weighted mixture of distillation objectives:

$$\min_{S, \mathbf{w}} \mathbb{E}_{x \sim \mathcal{D}_{\mathbf{w}}} [\ell(S(x), T(x))] \quad \text{s.t.} \quad \mathbf{w} \in \Delta_{|\mathcal{C}|}, \text{cost}(S; S_0) \leq B. \quad (1)$$




Figure 1: Exploratory study. (a–b) Cross-capability transfer under single-capability distillation: larger budgets sharpen target gains but expose stronger negative transfer. (c–d) Budget waste: extra tokens yield smaller target gains while increasing harmful spillover to non-target capabilities.

where $\mathcal{D}_{\mathbf{w}} := \sum_{c \in \mathcal{C}} w_c \mathcal{D}_c$, and $\text{cost}(S; S_0)$ measures the distillation token budget consumed to obtain S from S_0 . Here $\mathbf{w} = (w_c)_{c \in \mathcal{C}}$ lies in the probability simplex $\Delta_{|\mathcal{C}|} = \{\mathbf{w} \in \mathbb{R}^{|\mathcal{C}|} \mid w_c \geq 0, \sum_{c \in \mathcal{C}} w_c = 1\}$, which specifies the allocation of training effort across capabilities, and ℓ denotes a standard distillation loss, such as token-level cross-entropy or logit-matching between the student’s and the teacher’s output distributions.

We therefore conduct an empirical study to identify and quantify how allocating a limited distillation budget to specific capabilities induces systematic performance changes in other capabilities, thereby exposing the intrinsic interdependence structure among capabilities in LLMs.

Definition 2.2 (Capability interdependence). Let S_0 be an initial student model, and let $S(\mathbf{w}, B)$ denote the student obtained after capability distillation under allocation $\mathbf{w} \in \Delta_{|\mathcal{C}|}$ and budget B . For each capability c_i , define the one-hot allocation vector $\mathbf{w}^{(i)} \in \Delta_{|\mathcal{C}|}$ by $w_i^{(i)} = 1$ and $w_j^{(i)} = 0$ for all $j \neq i$. The *capability interdependence* under capability distillation is characterized by the capability transfer matrix

$$\mathbf{T}_{ij}(B) := s_j(S(\mathbf{w}^{(i)}, B)) - s_j(S_0),$$

where $\mathbf{T}(B) \in \mathbb{R}^{|\mathcal{C}| \times |\mathcal{C}|}$, $\mathbf{T}_{ij}(B)$ quantifies the effect of distilling toward capability c_i on the performance of capability c_j . Non-zero off-diagonal entries indicate interdependence between capabilities under shared representations.

2.3 EXPLORATORY STUDY AND PROBLEM FORMULATION

We distill Llama-3.3-70B-Instruct into Llama-3.1-8B-Instruct under fixed token budgets, one target capability at a time, and evaluate the student on all benchmarks to form the transfer matrix of Definition 2.2 (strategies and settings in Appendix B). All benchmark metrics are accuracies or success rates on a 0–100 scale, a capability score is the unweighted mean over its benchmarks, and every entry of the transfer matrix is an absolute difference in these points (Definition 2.2); no further rescaling is applied (Appendix B).

Observation 1: Distilling a specific capability redistributes performance across other capabilities in a budget-dependent manner. Figure 1(a–b) reports the visualization of the capability transfer matrices under different token budgets. Here, the off-diagonal mass is consistently non-zero, meaning optimizing a single capability alters other capabilities rather than leaving them unchanged. Besides, the structure varies with budget: at small budgets, changes are generally weak and diffuse, whereas at large budgets the diagonal strengthens, but negative off-diagonals become more visible for specific non-targets. Together, these results indicate that capability distillation induces a structured, **budget-dependent** redistribution of performance across capabilities.

Observation 2: Capability distillation exhibits substantial budget waste due to diminishing returns and negative spillover. Under a fixed distillation budget, additional tokens for a target capability should ideally provide meaningful target improvement without degrading other capabilities needed downstream. We test this using two diagnostics derived from the capability transfer matrix $\mathbf{T}(B)$. First, *diminishing returns* measures whether marginal target gains shrink as the budget

increases, comparing 20M→80M against 80M→150M for the same target. Second, *negative spillover* measures the drop on non-target capabilities as more budget is assigned to target.

Figure 1c shows that, for the strongest-improving targets, the 20M→80M gain is consistently larger than the 80M→150M gain, indicating target-capability saturation. Figure 1d further shows that collateral harm to non-target capabilities grows with budget. Together, these trends reveal budget waste: beyond a capability-specific knee point, extra tokens yield limited target benefit while amplifying cross-capability degradation, motivating allocation strategies that explicitly trade off target gains against spillover. **Neither observation is new in isolation: off-target drops are a form of catastrophic forgetting (Kirkpatrick et al., 2017; Luo et al., 2023), and saturating target gains are ordinary diminishing returns. What allocation can exploit is their *structure*: transfer is signed and source-specific, it changes with budget, and the two effects interact, so a token’s marginal value depends on where the budget has already gone.**

Overall, these observations suggest that effective capability distillation under a fixed budgets requires jointly reasoning about task-relevant capabilities, cross-capability interactions, and budget allocation. We formalize this as a budget allocation problem across interdependent capabilities:

Problem 2.3 (Budget allocation for interdependent capability distillation). *Given a teacher model T , an initial student model S_0 , a downstream task τ , and a total distillation budget B , consider a sequential distillation process over steps $t = 1, \dots, T$ with cumulative cost not exceeding B . At each step t , a capability allocation vector $\mathbf{w}_t \in \Delta_{|C|}$ is selected and applied to update the student, yielding S_t . Let $U_\tau(\cdot)$ is the task-dependent utility function, our target is to determine an allocation policy $\{\mathbf{w}_t\}_{t=1}^T$ that maximizes the performance of the final student model:*

$$\max_{\{\mathbf{w}_t\}_{t=1}^T} U_\tau(S_T), \text{ s.t. } \sum_{t=1}^T \text{cost}(S_{t-1} \rightarrow S_t) \leq B \quad (2)$$

3 METHODOLOGY

ReAD treats Problem 2.3 as state-dependent budget allocation over capabilities. At each step, it estimates task requirements, generates capability-targeted teacher supervision, and uses an uncertainty-aware contextual bandit to choose the next token allocation. As the student profile changes, ReAD can shift budget away from saturated or high-spillover capabilities. **ReAD is off-policy in the sense of Agarwal et al. (2024): the student trains on teacher completions to generated prompts and never samples its own training sequences. The allocator is orthogonal to the inner objective, and on-policy losses (Gu et al., 2024; Ko et al., 2024) could replace the token-level loss in Section 3.3 unchanged.**

3.1 IDENTIFYING TASK-ESSENTIAL CAPABILITIES

Task requirement vector. For each downstream task τ , ReAD constructs a task card $\mathcal{D}_\tau^{\text{spec}}$ from non-test data, containing a short task description, input/output format, evaluation target, and a few representative exemplars, with the test split reserved only for final reporting. From this task card, ReAD estimates a requirement vector $\mathbf{r}_\tau \in \Delta_{|C|}$, where $r_{\tau,c}$ indicates how useful capability c is for improving the task utility $U_\tau(\cdot)$. We use $\mathbf{r}_\tau$ as an actionable allocation prior: high-mass capabilities are prioritized for improvement, and drops on these capabilities are treated as harmful spillover. Since token allocations cannot be negative, $\mathbf{r}_\tau$ is constrained to the simplex; negative sensitivities are captured instead by signed capability changes and the spillover penalty in Section 3.3.

Requirement identifier. We learn a lightweight identifier g_ϕ that maps the task card to the requirement vector $\mathbf{r}_\tau = g_\phi(\mathcal{D}_\tau^{\text{spec}}) \in \Delta_{|C|}$. The identifier is a small Transformer encoder over the task card, followed by a two-layer MLP and a softmax head. Since $\mathbf{r}_\tau$ is not directly observed, we supervise it through low-budget interventions in capability space.

Local supervision signal. Let $\mathbf{s}(S) \in \mathbb{R}^{|C|}$ denote the measured capability profile of a student S , and write $U_\tau(S) = F_\tau(\mathbf{s}(S))$. For a small intervention that changes the profile from $\mathbf{s}$ to $\mathbf{s} + \Delta\mathbf{s}$,

$$F_\tau(\mathbf{s} + \Delta\mathbf{s}) = F_\tau(\mathbf{s}) + \nabla_{\mathbf{s}} F_\tau(\mathbf{s})^\top \Delta\mathbf{s} + o(\|\Delta\mathbf{s}\|), \quad (3)$$

so the utility change is locally approximated by

$$\Delta U_\tau = U_\tau(S') - U_\tau(S) \approx \nabla_s F_\tau(\mathbf{s})^\top \Delta \mathbf{s}. \quad (4)$$

Thus, $\mathbf{r}_\tau$ is trained as a simplex-constrained surrogate for the beneficial part of the local utility sensitivity, e.g., $\mathbf{r}_\tau \approx \Pi_\Delta([\nabla_s F_\tau(\mathbf{s})]_+)$.

Offline pretraining. We construct an intervention set $\mathcal{Z} = \{(\mathcal{D}_\tau^{\text{spec}}, \Delta \mathbf{s}_\tau^{(m)}, \Delta U_\tau^{(m)})\}_{\tau, m}$ from auxiliary tasks and non-test splits. For each task τ , we sample sparse allocation vectors $\mathbf{w}^{(m)} \in \Delta_{|\mathcal{C}|}$ by choosing a support set $\mathcal{I}^{(m)} \subseteq \mathcal{C}$, sampling nonzero weights from a Dirichlet distribution, and renormalizing them. A short probe distillation under a small budget b_{probe} produces an intervened student $S^{(m)}$, from which we measure

$$\Delta U_\tau^{(m)} := U_\tau(S^{(m)}) - U_\tau(S_0), \quad \Delta \mathbf{s}_\tau^{(m)} := \mathbf{s}(S^{(m)}) - \mathbf{s}(S_0). \quad (5)$$

We pretrain g_ϕ by predicting the observed utility change from the induced capability change:

$$\min_{\phi} \sum_{(\tau, m)} \left(\Delta U_\tau^{(m)} - g_\phi(\mathcal{D}_\tau^{\text{spec}})^\top \Delta \mathbf{s}_\tau^{(m)} \right)^2 + \lambda_{\text{ent}} \mathcal{H}(g_\phi(\mathcal{D}_\tau^{\text{spec}})). \quad (6)$$

The positive entropy penalty discourages near-uniform requirement vectors and makes the allocation prior more discriminative. At deployment time, ReAD computes $\mathbf{r}_\tau$ once and reuses it throughout the distillation run. Since the identifier is trained once and reused across tasks, it is not counted as part of the per-run student distillation budget.

3.2 ON-THE-FLY CAPABILITY-TARGETED DATA GENERATION

Capability-conditioned templates. Given $\mathbf{r}_\tau$, ReAD generates training examples on the fly from a capability-labeled template library. For each capability $c \in \mathcal{C}$, a template set $\mathcal{P}_c$ specifies an instruction scaffold, typed slots, and output-format constraints. Slot values are filled using deterministic seeds. The generator controls the form of teacher supervision

Difficulty scoring. For prompt x in capability c , we compute a deterministic difficulty score

$$d_c(x) = \frac{1}{|\mathcal{F}_c|} \sum_{f \in \mathcal{F}_c} \frac{f(x) - \min_{x' \in \mathcal{Q}_c} f(x')}{\max_{x' \in \mathcal{Q}_c} f(x') - \min_{x' \in \mathcal{Q}_c} f(x') + \epsilon}, \quad (7)$$

where $\mathcal{F}_c$ is the set of capability-specific control factors and $\mathcal{Q}_c$ is a calibration prompt pool for capability c . For example, reasoning templates use factors such as the number of constraints, reasoning hops, and symbolic complexity; code templates use the number of required functions and tests; steerability templates use the number of rules and schema fields. We split each capability’s calibration pool into easy, medium, and hard buckets by tertiles of $d_c(x)$.

Curriculum and sampling. At decision step t , ReAD samples a capability label according to the selected allocation $\mathbf{w}_t$, instantiates a prompt from the corresponding template set, queries the teacher T for a completion y , and immediately distills on (x, y) . The curriculum only adjusts the within-capability difficulty mix: early steps emphasize easy and medium examples, while later steps increase the probability of hard examples. It does not choose the capability allocation. We also track recent template frequencies and cap each template’s share so that no single scaffold dominates.

3.3 CONTEXTUAL BANDIT FOR CAPABILITY ALLOCATION

Decision loop. ReAD runs for $T_{\text{step}} = 20$ decision steps, **with the bandit update, data resampling, and proxy refresh at every step (Section 5)**. At step t , the context summarizes the task requirement, current student profile, remaining budget, and recent allocation history:

$$\mathbf{x}_t = [\mathbf{r}_\tau; \mathbf{s}^{\text{probe}}(S_t); b_t; \boldsymbol{\rho}_t], \quad (8)$$

where $\boldsymbol{\rho}_t$ stores a short record of recent allocations and observed gains. The profile $\mathbf{s}^{\text{probe}}(S_t)$ is computed with a small fixed monitoring suite, while full held-out benchmark evaluation is used only for final reporting and coarse calibration.

Candidate allocations and distillation. At each step, ReAD selects an allocation $\mathbf{w}_t$ from a finite candidate set $\mathcal{A}(\tau)$ consisting of local perturbations around the previous allocation and sparse actions concentrated on the top- k capabilities under $\mathbf{r}_\tau$. Weights are chosen from a fixed grid and renormalized to sum to one. Given $\mathbf{w}_t$, ReAD samples capability-targeted data and updates the student by minimizing $\mathcal{L}_{\text{distill}}(\theta_t) = -\mathbb{E}_{(x,y)} \sum_{j=1}^{|y|} \log p_{\theta_t}(y_j \mid x, y_{<j})$ to produce S_{t+1} .

Proxy reward. After the update, ReAD measures the probe-profile change $\Delta \mathbf{s}_t^{\text{probe}} = \mathbf{s}^{\text{probe}}(S_{t+1}) - \mathbf{s}^{\text{probe}}(S_t)$ and forms

$$\hat{R}_t = \mathbf{r}_\tau^\top \Delta \mathbf{s}_t^{\text{probe}} - \beta \text{Spill}_t - \lambda \text{cost}_t, \quad \text{Spill}_t = \sum_{c \in \mathcal{C}_{\text{ess}}} r_{\tau,c} [-\Delta s_{t,c}^{\text{probe}}]_+ \quad (9)$$

Here $\mathcal{C}_{\text{ess}}$ contains the top- k capabilities under $\mathbf{r}_\tau$, and cost_t is the token budget consumed at the step. The first term rewards task-aligned capability gains, the second penalizes regressions on task-essential capabilities, and the third enforces budget awareness.

Reward model and UCB allocation. We append $(\mathbf{x}_t, \mathbf{w}_t, \hat{R}_t)$ to a history buffer and train an ensemble of J two-layer MLP reward regressors $\{h_{\eta_j}\}_{j=1}^J$. Each regressor maps $(\mathbf{x}, \mathbf{w})$ to a scalar next-step reward and is trained on a bootstrap resample of the 10% profiling split plus logged transitions from earlier steps, yielding about 3.2k–4.9k examples per task in our experiments. For a candidate allocation $\mathbf{w}$, the ensemble mean and uncertainty are

$$\mu(\mathbf{x}_t, \mathbf{w}) = \frac{1}{J} \sum_{j=1}^J h_{\eta_j}(\mathbf{x}_t, \mathbf{w}), \quad \sigma(\mathbf{x}_t, \mathbf{w}) = \sqrt{\frac{1}{J-1} \sum_{j=1}^J (h_{\eta_j}(\mathbf{x}_t, \mathbf{w}) - \mu(\mathbf{x}_t, \mathbf{w}))^2}. \quad (10)$$

ReAD then selects the next allocation using an upper-confidence rule,

$$\mathbf{w}_{t+1} = \arg \max_{\mathbf{w} \in \mathcal{A}(\tau)} \mu(\mathbf{x}_{t+1}, \mathbf{w}) + \kappa \sigma(\mathbf{x}_{t+1}, \mathbf{w}), \quad (11)$$

with $b_{t+1} = b_t - \text{cost}_t$. At coarse development checkpoints, we fit an affine calibration from cumulative proxy reward to observed utility change and reduce exploration if the development utility stagnates. Unlike static mixture-regression approaches such as RegMix (Liu et al., 2024), ReAD repeatedly re-estimates the best next allocation from the current student state and remaining budget.

4 THEORETICAL ANALYSIS

This section provides a local theoretical account of ReAD’s allocation problem. We characterize cross-capability transfer under shared representations, show when diminishing returns make continued allocation wasteful, and motivate uncertainty-aware allocation under noisy reward estimates.

4.1 LOCAL CROSS-CAPABILITY TRANSFER

Let $\theta_t \in \mathbb{R}^p$ denote the parameters of the student S_t , and let $\Delta \theta_t = \theta_t - \theta_{t-1}$. We analyze one decision step, which is the regime in which ReAD re-estimates its allocation.

Assumption 4.1 (Mixture-structured update). At interval t , given an allocation vector $\mathbf{w}_t \in \Delta_{|\mathcal{C}|}$, the expected update admits a local mixture approximation $\mathbb{E}[\Delta \theta_t \mid \mathbf{w}_t, \theta_{t-1}] \approx \sum_{c \in \mathcal{C}} w_{t,c} d_c(\theta_{t-1})$, where $d_c(\theta)$ is the capability-conditioned local update direction.

Assumption 4.2 (Local capability readout). The measured capability profile $\mathbf{s}(S)$ is represented locally as a differentiable function $\tilde{\mathbf{s}}(\theta)$. Near θ_{t-1} ,

$$\Delta s_{t,c} = [\mathbf{s}(S_t)]_c - [\mathbf{s}(S_{t-1})]_c \approx \nabla_\theta [\tilde{\mathbf{s}}(\theta_{t-1})]_c^\top \Delta \theta_t. \quad (12)$$

Proposition 4.3 (Local transfer decomposition). Under the two local approximations above, the expected change in capability c satisfies

$$\Delta s_{t,c} \approx \sum_{c' \in \mathcal{C}} w_{t,c'} \Gamma_{c,c'}(\theta_{t-1}), \quad \Gamma_{c,c'}(\theta) := \nabla_\theta [\tilde{\mathbf{s}}(\theta)]_c^\top d_{c'}(\theta). \quad (13)$$

Interpretation. Γ is the local analogue of the empirical transfer matrix: diagonal terms are intended gains, and off-diagonal terms are nonzero whenever the update direction for c' is not orthogonal to the readout direction of c , which is expected under shared representations. This is why allocating budget to one capability alters others at the decision-step level.

4.2 BUDGET WASTE FROM DIMINISHING RETURNS

The previous subsection explains spillover. We now isolate diminishing returns with a simpler cumulative-budget proxy. Let $b_c \geq 0$ be the total number of distillation tokens assigned to capability c , with $\sum_c b_c \leq B$, and let $\mathbf{r}_\tau \in \Delta_{|\mathcal{C}|}$ be the task requirement vector.

Assumption 4.4 (Concave task-aligned gains). For task τ and capability c , there is a nondecreasing concave gain function $G_{\tau,c} : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ such that $\mathbb{E}[U_\tau(S) - U_\tau(S_0)] \approx \sum_{c \in \mathcal{C}} r_{\tau,c} G_{\tau,c}(b_c)$ where $\sum_{c \in \mathcal{C}} b_c \leq B$.

Consider a restricted distillation strategy that allocates budget only to a subset $\mathcal{K} \subseteq \mathcal{C}$:

$$\max_{\{b_c \geq 0\}_{c \in \mathcal{K}}} \sum_{c \in \mathcal{K}} r_{\tau,c} G_{\tau,c}(b_c) \quad \text{s.t.} \quad \sum_{c \in \mathcal{K}} b_c \leq B. \quad (14)$$

Proposition 4.5 (Weighted marginal gain criterion). Any optimum $\{b_c^*\}_{c \in \mathcal{K}}$ of Eq. (14) equalizes weighted marginal gains among active capabilities: there exists $\nu^* \geq 0$ such that

$$r_{\tau,c} G'_{\tau,c}(b_c^*) \leq \nu^* \quad \forall c \in \mathcal{K}, \quad r_{\tau,c} G'_{\tau,c}(b_c^*) = \nu^* \quad \text{if } b_c^* > 0. \quad (15)$$

Moreover, in the single-capability case targeting $\bar{c}$, shifting an infinitesimal budget from $\bar{c}$ to another task-relevant capability c' improves the proxy objective whenever $r_{\tau,c'} G'_{\tau,c'}(0) > r_{\tau,\bar{c}} G'_{\tau,\bar{c}}(b_{\bar{c}})$.

Implication for ReAD. This formalizes budget waste: once the target saturates, further tokens have lower task-aligned marginal value than tokens assigned elsewhere. ReAD’s reward in Eq. (9) is richer than this proxy because it also penalizes observed spillover, so the allocator optimizes a step-level reward that accounts for both diminishing returns and cross-capability degradation.

4.3 WHY UNCERTAINTY-AWARE ALLOCATION HELPS

At step t , ReAD chooses from a finite action set $\mathcal{A}(\tau)$ using the ensemble mean $\mu(\mathbf{x}_t, \mathbf{w})$ and uncertainty $\sigma(\mathbf{x}_t, \mathbf{w})$ in Eq. (11). Let $\bar{R}(\mathbf{x}, \mathbf{w}) := \mathbb{E}[\hat{R}_t \mid \mathbf{x}_t = \mathbf{x}, \mathbf{w}_t = \mathbf{w}]$ be the conditional expected proxy reward of action $\mathbf{w}$ in context $\mathbf{x}$.

Assumption 4.6 (Calibrated ensemble uncertainty). For all encountered $(\mathbf{x}_t, \mathbf{w})$, with high probability, $|\mu(\mathbf{x}_t, \mathbf{w}) - \bar{R}(\mathbf{x}_t, \mathbf{w})| \leq \kappa \sigma(\mathbf{x}_t, \mathbf{w})$.

Proposition 4.7 (Step-level regret bound). Let $\mathbf{w}_t^* \in \arg \max_{\mathbf{w} \in \mathcal{A}(\tau)} \bar{R}(\mathbf{x}_t, \mathbf{w})$. Under Assumption 4.6, the action selected by Eq. (11) satisfies $\bar{R}(\mathbf{x}_t, \mathbf{w}_t^*) - \bar{R}(\mathbf{x}_t, \mathbf{w}_t) \leq 2\kappa \sigma(\mathbf{x}_t, \mathbf{w}_t)$.

Proof sketch. By calibration, the true reward of any action is at most its upper confidence score, and the selected action maximizes this score. Applying the same calibration inequality to the selected action yields the bound. This result does not imply a global optimum for nonconvex fine-tuning; it justifies the local ranking rule used by ReAD at each budget step.

5 EXPERIMENTAL SETTING

Models and budgets. We mainly distill Llama-3.3-70B-Instruct into Llama-3.1-8B-Instruct, and report Qwen2.5-72B-Instruct to Qwen2.5-14B-Instruct results in Appendix D. All methods use the same student checkpoint, training recipe, and 20M/150M distillation-token budgets.

Evaluation protocol. We evaluate eight capabilities using Appendix Table 3. For each benchmark, 10% of non-test data is used for profiling, requirement inference, and allocation updates; the held-out split is reserved for final reporting. Generation uses temperature 0 with the same prompt template across methods. We report mean $\pm$ standard deviation over three seeds.

Baselines. For each capability, we build a held-out-filtered distillation pool from its benchmark prompts. For multi-dataset capabilities, we use balanced dataset-level sampling. The six main baselines combine three teacher signals—final response (Resp), chain-of-thought plus response (CoT), and token-level logits (Logit)—with two objectives, SFT and KD. Each baseline is one-hot at the capability level, spending the full budget on the target-capability pool. We also compare against four allocation baselines that share ReAD’s generator, teacher, and budget and differ only in how

Table 1: Capability profile under a 20M-token budget on Llama. ReAD achieves the strongest average profile while using the same student initialization and token budget as all baselines.

Method	General	Steer.	Reason.	Math	Code	Tool	LCU	Multi.
Teacher only	48.13 $\pm$ 0.35	89.98 $\pm$ 0.40	10.51 $\pm$ 0.12	48.34 $\pm$ 0.55	88.40 $\pm$ 0.70	31.90 $\pm$ 0.45	36.20 $\pm$ 0.40	91.10 $\pm$ 0.55
Student only	31.09 $\pm$ 0.40	49.22 $\pm$ 0.55	8.72 $\pm$ 0.15	15.56 $\pm$ 0.60	72.60 $\pm$ 0.95	25.83 $\pm$ 0.60	30.40 $\pm$ 0.45	68.90 $\pm$ 0.80
Resp-SFT	33.28 $\pm$ 0.47	52.84 $\pm$ 0.71	9.05 $\pm$ 0.18	17.12 $\pm$ 0.83	74.06 $\pm$ 1.09	26.92 $\pm$ 0.78	31.28 $\pm$ 0.51	70.02 $\pm$ 0.93
CoT-SFT	33.02 $\pm$ 0.52	52.16 $\pm$ 0.76	9.46 $\pm$ 0.14	18.58 $\pm$ 0.77	73.82 $\pm$ 1.02	26.81 $\pm$ 0.74	31.19 $\pm$ 0.57	69.78 $\pm$ 0.88
Logit-SFT	33.61 $\pm$ 0.44	53.04 $\pm$ 0.69	9.12 $\pm$ 0.17	17.63 $\pm$ 0.79	74.48 $\pm$ 1.15	27.03 $\pm$ 0.70	31.42 $\pm$ 0.53	70.21 $\pm$ 0.90
Resp-KD	33.83 $\pm$ 0.50	53.62 $\pm$ 0.73	9.18 $\pm$ 0.16	17.84 $\pm$ 0.74	74.23 $\pm$ 1.07	27.18 $\pm$ 0.76	31.58 $\pm$ 0.55	70.48 $\pm$ 0.95
CoT-KD	33.46 $\pm$ 0.46	52.98 $\pm$ 0.78	9.62 $\pm$ 0.13	19.34 $\pm$ 0.72	74.03 $\pm$ 1.10	27.06 $\pm$ 0.72	31.47 $\pm$ 0.58	70.08 $\pm$ 0.89
Logit-KD	34.02 $\pm$ 0.43	54.21 $\pm$ 0.66	9.25 $\pm$ 0.15	18.20 $\pm$ 0.76	74.62 $\pm$ 1.12	27.41 $\pm$ 0.69	31.72 $\pm$ 0.50	70.81 $\pm$ 0.91
ReAD	35.02 $\pm$ 0.41	58.03 $\pm$ 0.61	10.01 $\pm$ 0.12	21.06 $\pm$ 0.68	77.06 $\pm$ 0.97	29.82 $\pm$ 0.64	33.58 $\pm$ 0.49	73.04 $\pm$ 0.84




Figure 2: Performance versus token budget for four bottleneck capabilities. Curves show mean $\pm$ standard deviation over three seeds.

$\mathbf{w}_t$ is chosen (Table 2): *Uniform static* ($\mathbf{w}_t = 1/|C|$), *Task-static mix* ($\mathbf{w}_t = \mathbf{r}_\tau$), *Greedy one-step* (re-selects $\mathbf{w}_t$ each step by the ensemble mean alone, $\kappa = 0$), and *Grid-searched* (best fixed allocation in $\mathcal{A}(\tau)$ on the profiling split). These are the matched reference points for adaptive allocation.

ReAD implementation. ReAD runs for 20 decision steps, with bandit updates, proxy refreshes, and data resampling every 1.0M tokens in the 20M setting and every 7.5M tokens in the 150M setting. Proxy refresh adds 3.12% wall-clock overhead, the full online loop adds 5.63%, and the reusable requirement identifier costs 3.79 GPU-hours to train once. **Cost accounting.** B counts the teacher-generated tokens the student is trained on and is identical for all methods. Two ReAD-specific costs fall outside B and are reported separately rather than folded into it: the identifier is trained once from auxiliary tasks and reused for every task and budget, so its 3.79 GPU-hours are amortized, and the fixed probe suite is the 3.12% overhead above.

6 EXPERIMENTS

We study four questions: **RQ1** utility gains under the same budget; **RQ2** reduced waste and spillover; **RQ3** the value of adaptive bandit scheduling over static or greedy alternatives; and **RQ4** transfer beyond capability-aligned benchmarks.

6.1 RESULTS AND ANALYSIS

RQ1: Utility under a fixed budget. Table 1 shows that ReAD achieves the strongest capability profile at 20M tokens, **outperforming the six one-hot baselines on every capability**. The gains are especially pronounced on bottleneck capabilities such as Steerability, Math, Tool Use, and LCU. **Since each one-hot baseline spends the whole budget on one capability, this isolates the value of distributing the budget at all; the matched comparison for how to distribute it is Table 2(b) (RQ3).** The 150M Llama results and Qwen results in Appendix D show the same ordering, suggesting that the improvement is not tied to a single teacher-student pair.

RQ2: Budget efficiency. Figure 2 reports budget-scaling curves for Reasoning, Math, Tool Use, and LCU. ReAD achieves higher performance at the same budget and maintains the strongest scaling trend, indicating that adaptive allocation reduces wasted tokens on low-marginal-return updates.

Table 2: **Budget-allocation diagnostics.** (a) 150M-token gain–spillover trade-off from the capability transfer matrix: ReAD keeps near-best on-target gain with much less off-target degradation. (b) Matched-budget allocation schedules under the same generator.

(a) Harmful transfer at 150M				(b) Allocation schedules		
Method	On-tgt $\uparrow$	Off-tgt $\Delta \uparrow$	Neg. $\downarrow$	Method	20M $\uparrow$	150M $\uparrow$
Single-cap. KD	4.79$\pm$0.29	-1.57 $\pm$ 0.21	0.43 $\pm$ 0.03	Uniform static	60.94 $\pm$ 0.48	64.73 $\pm$ 0.36
Task-static mix	4.61 $\pm$ 0.25	-0.87 $\pm$ 0.17	0.29 $\pm$ 0.03	Task-static mix	61.71 $\pm$ 0.43	65.52 $\pm$ 0.33
Greedy one-step	4.72 $\pm$ 0.23	-0.94 $\pm$ 0.19	0.31 $\pm$ 0.03	Greedy one-step	62.03 $\pm$ 0.41	66.01 $\pm$ 0.30
Grid-searched	4.66 $\pm$ 0.24	-0.71 $\pm$ 0.16	0.25 $\pm$ 0.03	Grid-searched	61.88 $\pm$ 0.42	65.81 $\pm$ 0.31
ReAD	4.79 $\pm$ 0.24	-0.38$\pm$0.12	0.17$\pm$0.02	ReAD	63.04$\pm$0.39	67.21$\pm$0.28




Figure 3: Ablation of ReAD components: average benchmark score over the eight capabilities at 20M and 150M tokens (higher is better).

RQ2: Gain–spillover trade-off. Table 2(a) shows that ReAD improves the gain–spillover trade-off. Single-capability KD achieves strong target gains but causes larger off-target degradation, while static mixing reduces spillover at the cost of weaker target improvement. ReAD better balances the two: it keeps strong on-target gains while suppressing negative transfer across non-target capabilities.

RQ3: Adaptive scheduling. Table 2(b) compares ReAD with stronger allocation baselines under the same generator and token budgets. Static schedules cannot adapt once marginal returns and spillover change during training, while the greedy heuristic ignores uncertainty and remaining budget. ReAD outperforms both, with a larger margin at 150M where saturation and spillover are stronger. These margins are smaller than those over one-hot baselines (+1.0/+1.2 over Greedy one-step and +1.2/+1.4 over Grid-searched at 20M/150M, with seed standard deviations of 0.3–0.4; Welch tests in Appendix Table 7), as every allocation baseline already mixes capabilities; the remaining gain is the value of re-estimating the mix as the student and the remaining budget change.

Component ablation. We remove one component at a time with everything else fixed: without the requirement identifier the allocation is uniform, without adaptation it is fixed, and without interaction awareness the spillover penalty is dropped. Figure 3 shows that identification and adaptation drive the largest gains, while interaction modeling adds a further improvement by reducing harmful transfer.

RQ4: Transfer beyond aligned benchmarks. On held-out XSTest (Röttger et al., 2024), with the same identifier and allocation pipeline and no new capability head, ReAD improves both safe- and unsafe-refusal metrics over the best SFT/KD baseline (Appendix Table 9).

7 CONCLUSIONS

We study capability distillation under a fixed token budget and find systematic cross-transfer and diminishing returns, often with increasing harmful spillover on task-critical behaviors. We propose ReAD, a reinforcement-guided framework that identifies task-essential capabilities and adaptively allocates budget across interdependent capabilities via on-the-fly supervision and an uncertainty-aware contextual bandit. Experiments show ReAD improves downstream utility under the same budget while reducing **wasted distillation effort**.

AI USE STATEMENT

In this work, we used generative AI tools to assist with implementing parts of the experimental code (e.g., data processing, evaluation, and plotting scripts) and to improve the writing of the manuscript (grammar, wording, and formatting). We have not used generative AI tools to generate synthetic datasets, develop the theoretical model or conceptual framework, formulate or prove mathematical claims, propose or refine hypotheses, design the research methodology or experiments, or interpret results; the remaining tasks with required disclosure (e.g., translation and qualitative or thematic data analysis) are not applicable to this work. We have reviewed all AI-assisted work: AI-assisted code was checked and tested by the authors before use, and all AI-edited text was verified by the authors against the intended content. We take responsibility for the final content of this work, including text, claims, and artifacts produced with the aid of generative AI.

ETHICS STATEMENT

This work does not involve human subjects, crowdsourcing, or personally identifiable data. All experiments use publicly available pretrained models and public benchmarks, and evaluation is fully automated. ReAD is a method for allocating a fixed distillation budget across capabilities; it does not introduce new datasets or release new model weights. As with any distillation method, distilled students may inherit biases or unsafe behaviors from the teacher and the training data, and more efficient distillation may lower the cost of producing specialized models with dual-use capabilities. We discuss these risks and the corresponding mitigations (task-specific safety evaluation, bias and robustness audits, and monitoring of off-target capability changes) in Appendix F.

REPRODUCIBILITY STATEMENT

We have made efforts to ensure that our results are reproducible. Section 5 specifies the teacher-student model pairs, token budgets, evaluation protocol, baselines, and the ReAD update schedule, and Section 3 describes each component of ReAD. Appendix A lists the benchmarks, metrics, and data splits used for each capability, Appendix B details the setting of the exploratory study, and Appendix D reports additional experiments and ablations. The assumptions and full derivations for the theoretical results in Section 4 are given in Appendix C. All reported numbers are means and standard deviations over three seeds. An anonymized repository with our code, configuration files, and instructions for reproducing the main results is available at <https://anonymous.4open.science/r/ReAD-62B7> (anonymized for review).

REFERENCES

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations (ICLR)*, 2024.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. LongBench: A bilingual, multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3119–3137, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.172. URL <https://aclanthology.org/2024.acl-long.172>.
- Zinuo Cai, Rongbo Ma, Yicheng Fu, Weishan Zhang, Ruhui Ma, and Haibing Guan. Llmaas: Serving large language models on trusted serverless computing platforms. *IEEE Transactions on Artificial Intelligence*, 2024.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan,

- Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code. 2021.
- Daixuan Cheng, Shaohan Huang, and Furu Wei. Adapting large language models via reading comprehension. In *The Twelfth International Conference on Learning Representations*, 2023.
- Wei-Lin Chiang, Zhuohan Li, Ziqing Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E Gonzalez, et al. Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality. See <https://vicuna.lmsys.org> (accessed 14 April 2023), 2(3):6, 2023.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. *ArXiv*, abs/1803.05457, 2018.
- Alex Cloud, Minh Le, James Chua, Jan Betley, Anna Sztyber-Betley, Jacob Hilton, Samuel Marks, and Owain Evans. Subliminal learning: Language models transmit behavioral traits via hidden signals in data. *arXiv preprint arXiv:2507.14805*, 2025.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021.
- Luyang Fang, Xiaowei Yu, Jiazhang Cai, Yongkai Chen, Shushan Wu, Zhengliang Liu, Zhenyuan Yang, Haoran Lu, Xilin Gong, Yufang Liu, Terry Ma, Wei Ruan, Ali Abbasi, Jing Zhang, Tao Wang, Ehsan Latif, Wei Liu, Wei Zhang, Soheil Kolouri, Xiaoming Zhai, Dajiang Zhu, Wenxuan Zhong, Tianming Liu, and Ping Ma. Knowledge distillation and dataset distillation of large language models: Emerging trends, challenges, and future directions. *arXiv preprint arXiv:2504.14772*, 2025.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *NeurIPS*, 2021.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Cheng-Yu Hsieh, Chun-Liang Li, Chih-Kuan Yeh, Hootan Nakhost, Yasuhisa Fujii, Alexander Ratner, Ranjay Krishna, and Chen-Yu Lee. Distilling step-by-step! outperforming larger language models with less training data and smaller model sizes. *arXiv preprint arXiv:2305.02301*, 2023.
- Xiaoqi Jiao, Yichun Yin, Lifeng Shang, Xin Jiang, Xiao Chen, Linlin Li, Fang Wang, and Qun Liu. Tinybert: Distilling bert for natural language understanding. In *Findings of the association for computational linguistics: EMNLP 2020*, pp. 4163–4174, 2020.
- Yimin Jing, Renren Jin, Jiahao Hu, Huishi Qiu, Xiaohua Wang, Peng Wang, and Deyi Xiong. Followeval: A multi-dimensional benchmark for assessing the instruction-following capability of large language models. *arXiv preprint arXiv:2311.09829*, 2023.
- Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1317–1327, 2016.

- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards streamlined distillation for large language models. In *International Conference on Machine Learning (ICML)*, 2024.
- Lei Li, Yankai Lin, Shuhuai Ren, Peng Li, Jie Zhou, and Xu Sun. Dynamic knowledge distillation for pre-trained language models. 2021.
- Raymond Li, Erik Nijkamp, Swaroop Mishra, et al. Starcoder: May the source be with you! *arXiv preprint arXiv:2305.06161*, 2023.
- Chen Liang, Simiao Zuo, Qingru Zhang, Pengcheng He, Weizhu Chen, and Tuo Zhao. Less is more: Task-aware layer-wise distillation for language model compression. In *International Conference on Machine Learning*, pp. 20852–20867. PMLR, 2023.
- Qian Liu, Xiaosen Zheng, Niklas Muennighoff, Guangtao Zeng, Longxu Dou, Tianyu Pang, Jing Jiang, and Min Lin. Regmix: Data mixture as regression for language model pre-training. *arXiv preprint arXiv:2407.01492*, 2024.
- Yun Luo, Zhen Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *arXiv preprint arXiv:2308.08747*, 2023.
- Lucie Charlotte Magister, Jonathan Mallinson, Jakub Adamek, Eric Malmi, and Aliaksei Severyn. Teaching small language models to reason. In *Proceedings of the 61st annual meeting of the association for computational linguistics (volume 2: short papers)*, pp. 1773–1781, 2023.
- Baolin Peng, Chunyuan Li, Pengcheng He, Michel Galley, and Jianfeng Gao. Instruction tuning with gpt-4. *arXiv preprint arXiv:2304.03277*, 2023.
- Edoardo M. Ponti, Goran Glavaš, Olga Majewska, Qianchu Liu, Ivan Vulić, and Anna Korhonen. XCOPA: A multilingual dataset for causal commonsense reasoning. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020. URL <https://ducdauge.github.io/files/xcopa.pdf>.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*, 2024.
- Paul Röttger, Hannah Rose Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics*, 2024.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Freda Shi, Mirac Suzgun, Markus Freitag, Xuezhi Wang, Suraj Srivats, Soroush Vosoughi, Hyung Won Chung, Yi Tay, Sebastian Ruder, Denny Zhou, Dipanjan Das, and Jason Wei. Language models are multilingual chain-of-thought reasoners, 2022.
- Kumar Shridhar, Alessandro Stolfo, and Mrinmaya Sachan. Distilling reasoning capabilities into smaller language models. *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 7059–7073, 2023.
- Zhihong Sun, Chen Lyu, Bolun Li, Yao Wan, Hongyu Zhang, Ge Li, and Zhi Jin. Enhancing code generation performance of smaller models by distilling the reasoning ability of llms. *arXiv preprint arXiv:2403.13271*, 2024.

- Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V Le, Ed H Chi, Denny Zhou, , and Jason Wei. Challenging big-bench tasks and whether chain-of-thought can solve them. *arXiv preprint arXiv:2210.09261*, 2022.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B Hashimoto. Stanford alpaca: An instruction-following llama model, 2023.
- Wenhui Wang, Furu Wei, Li Dong, Hangbo Bao, Nan Yang, and Ming Zhou. Minilm: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. volume 33, pp. 5776–5788, 2020.
- Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhui Chen. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark, 2024. URL <https://arxiv.org/abs/2406.01574>.
- Sang Michael Xie, Hieu Pham, Xuanyi Dong, Nan Du, Hanxiao Liu, Yifeng Lu, Percy Liang, Quoc V. Le, Tengyu Ma, and Adams Wei Yu. DoReMi: Optimizing data mixtures speeds up language model pretraining. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, and Daxin Jiang. Wizardlm: Empowering large language models to follow complex instructions. *arXiv preprint arXiv:2304.12244*, 2023a.
- Qiantong Xu, Fenglu Hong, Bo Li, Changran Hu, Zhengyu Chen, and Jian Zhang. On the tool manipulation capability of open-source large language models, 2023b.
- Xiaohan Xu, Ming Li, Chongyang Tao, Tao Shen, Reynold Cheng, Jinyang Li, Can Xu, Dacheng Tao, and Tianyi Zhou. A survey on knowledge distillation of large language models. *arXiv preprint arXiv:2402.13116*, 2024.
- Fanjia Yan, Huanzhi Mao, Charlie Cheng-Jie Ji, Tianjun Zhang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Berkeley function calling leaderboard. https://gorilla.cs.berkeley.edu/blogs/8_berkeley_function_calling_leaderboard.html, 2024.
- Chuanpeng Yang, Yao Zhu, Wang Lu, Yidong Wang, Qian Chen, Chenlong Gao, Bingjie Yan, and Yiqiang Chen. Survey on knowledge distillation for large language models: methods, evaluation, and application. *ACM Transactions on Intelligent Systems and Technology*, 2024.
- Yuanhao Yue, Chengyu Wang, Jun Huang, and Peng Wang. Distilling instruction-following abilities of large language models with task-aware curriculum planning. *arXiv preprint arXiv:2405.13448*, 2024.
- Junjie Zhang, Ruobing Xie, Yupeng Hou, Xin Zhao, Leyu Lin, and Ji-Rong Wen. Recommendation as instruction following: A large language model empowered recommendation approach. *ACM Transactions on Information Systems*, 43(5):1–37, 2025.
- Yichi Zhang, Zhuo Chen, Yin Fang, Yanxi Lu, Li Fangming, Wen Zhang, and Huajun Chen. Knowledgeable preference alignment for llms in domain-specific question answering. In *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 891–904, 2024.
- Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, et al. A survey of large language models. *arXiv preprint arXiv:2303.18223*, 1(2), 2023.
- Qihuang Zhong, Liang Ding, Li Shen, Juhua Liu, Bo Du, and Dacheng Tao. Revisiting knowledge distillation for autoregressive language models. *arXiv preprint arXiv:2402.11890*, 2024.
- Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023.
- Xunyu Zhu, Jian Li, Yong Liu, Can Ma, and Weiping Wang. Distilling mathematical reasoning capabilities into small language models. *Neural Networks*, 179:106594, 2024.

TECHNICAL APPENDICES AND SUPPLEMENTARY MATERIAL

A BENCHMARK DATASETS

Table 3 summarizes the benchmark suite used to evaluate each capability. For capabilities associated with multiple benchmarks, we report the average score across the listed benchmarks.

Table 3: Benchmark suite for evaluating LLM capabilities.

Capability	Benchmarks	Metric
General	MMLU-Pro, BIG	Accuracy
Steerability	IFEval, FollowEval	Instruction-following accuracy
Reasoning	GPQA, ARC-Challenge	Normalized accuracy
Math	MATH Hard, GSM8K	Exact match
Code	HumanEval, MBPP	Pass@1
Tool Use	BFCL V4, ToolBench	Tool-call success rate
LCU	LongBench	Average normalized score
Multilingual	XCOPA, MGSM	Macro accuracy

- **MMLU-Pro** (Wang et al., 2024) tests broad knowledge and problem-solving across many academic subjects using multiple-choice questions.
- **BIG-bench (BIG)** (Suzgun et al., 2022) is a large collection of diverse tasks designed to probe general reasoning and knowledge beyond standard exams.
- **IFEval** (Zhou et al., 2023) measures instruction following by checking whether model outputs satisfy explicit, verifiable constraints in the prompt.
- **FollowEval** (Jing et al., 2023) evaluates how well a model follows user instructions under varied formats and constraint types, complementing IFEval with broader instruction patterns.
- **GPQA** (Rein et al., 2024) is a graduate-level question answering benchmark that stresses hard scientific reasoning with high-quality, expert-written questions.
- **ARC-Challenge** (Clark et al., 2018) contains difficult grade-school science questions that require multi-step reasoning and commonsense beyond keyword matching.
- **MATH Hard** (Hendrycks et al., 2021) is a challenging subset of competition-style math problems and is scored by exact final-answer match.
- **GSM8K** (Cobbe et al., 2021) focuses on grade-school math word problems that require intermediate reasoning steps, also scored by exact final-answer match.
- **HumanEval** (Chen et al., 2021) evaluates code generation via functional correctness on hand-written programming problems, reported as Pass@1.
- **MBPP** (Austin et al., 2021) is a Python programming benchmark with short programming tasks and unit tests, also evaluated by Pass@1.
- **BFCL V4** (Yan et al., 2024) measures tool-use ability by checking whether the model can produce correct function calls under realistic tool specifications.
- **ToolBench** (Xu et al., 2023b) evaluates tool use in more complex, multi-tool scenarios and is scored by tool-call success rate.
- **LongBench** (Bai et al., 2024) evaluates long-context understanding across multiple long-input tasks and reports an average normalized score.
- **XCOPA** (Ponti et al., 2020) tests multilingual causal commonsense reasoning across languages and is reported using macro accuracy.
- **MGSM** (Shi et al., 2022) is a multilingual version of grade-school math and is evaluated by macro accuracy across languages.

B EMPIRICAL STUDY SETTING

We study capability interactions under a fixed distillation budget by distilling a student model toward one target capability at a time and then evaluating the resulting student on *all* capabilities in Table 3. For each capability, we build a capability-specific distillation pool from the corresponding benchmark prompts, excluding any held-out evaluation examples to prevent leakage. We repeat this process under multiple token budgets (e.g., 20M, 80M, 150M tokens) to form a budget-dependent capability transfer matrix, where each entry is the score change on a destination capability when distilling toward a source capability at budget B . We use deterministic decoding for generation benchmarks (temperature = 0) with a fixed prompt template across methods, and we compute each capability score as the average over its associated benchmarks in the table. We instantiate a simple but representative set of capability distillation baselines by combining three supervision forms that expose different levels of teacher knowledge—*Resp* (final answer only), *CoT* (rationale plus final answer), and *Logit* (token-level soft targets)—with two standard objectives, *SFT* and *KD*, yielding six baselines in total.

Scoring and normalization. Each benchmark in Table 3 is scored with its standard metric (accuracy, exact match, pass@1, or tool-call success rate) expressed in percentage points; LongBench reports its own per-task normalized score on the same 0–100 scale. A capability score is the unweighted mean of its benchmark scores, and a transfer-matrix entry $\mathbf{T}_{ij}(B)$ is the absolute difference in capability score between the distilled and the initial student. No min-max rescaling or gap-closure normalization is applied to the reported tables or transfer matrices, so a change of one point is one percentage point of the underlying metric, and small absolute changes remain small in the figures.

C ADDITIONAL THEORETICAL ANALYSIS

Uncertainty-aware UCB allocation. Here, we conduct an uncertainty-aware UCB allocation analysis showing that action-level regret is controlled by predictive uncertainty. At interval t , ReAD maintains a bandit context

$$\mathbf{x}_t := [\mathbf{r}_\tau; \mathbf{s}(S_t); b_t], \quad (16)$$

where $\mathbf{r}_\tau \in \Delta$ is the task requirement vector, $\mathbf{s}(S_t) \in \mathbb{R}^{|C|}$ is the probe-defined capability profile of the current student, and b_t is the remaining budget. Let $\mathcal{A}(\tau)$ be the discrete candidate action set of allocation vectors. Define the conditional expected proxy reward

$$\bar{R}(\mathbf{x}, \mathbf{w}) := \mathbb{E}[\hat{R}_t \mid \mathbf{x}_t = \mathbf{x}, \mathbf{w}_t = \mathbf{w}], \quad (17)$$

where $\hat{R}_t$ is the proxy reward used as shown in Eq. (9).

Assumption C.1 (Calibrated uncertainty). For a confidence level $\delta \in (0, 1)$, with probability at least $1 - \delta$, for all encountered $(\mathbf{x}_t, \mathbf{w})$,

$$|\mu(\mathbf{x}_t, \mathbf{w}) - \bar{R}(\mathbf{x}_t, \mathbf{w})| \leq \sigma(\mathbf{x}_t, \mathbf{w}), \quad (18)$$

where μ and σ are the ensemble mean and standard deviation defined from the reward regressors.

Since ReAD selects allocations by the UCB rule in Eq. (11), we can compare its chosen action $\mathbf{w}_t$ to the optimal action $\mathbf{w}_t^*$ under the same context and bound the resulting gap in conditional expected reward using the calibrated uncertainty condition in Assumption C.1, which leads to the following instantaneous regret bound on the chosen action.

Proposition C.2 (Instantaneous regret bound). Let $\mathbf{w}_t^* \in \arg \max_{\mathbf{w} \in \mathcal{A}(\tau)} \bar{R}(\mathbf{x}_t, \mathbf{w})$. Under Assumption C.1,

$$\bar{R}(\mathbf{x}_t, \mathbf{w}_t^*) - \bar{R}(\mathbf{x}_t, \mathbf{w}_t) \leq 2\kappa \sigma(\mathbf{x}_t, \mathbf{w}_t), \quad (19)$$

and consequently,

$$\sum_{t=1}^T \left(\bar{R}(\mathbf{x}_t, \mathbf{w}_t^*) - \bar{R}(\mathbf{x}_t, \mathbf{w}_t) \right) \leq 2\kappa \sum_{t=1}^T \sigma(\mathbf{x}_t, \mathbf{w}_t). \quad (20)$$

Table 4: Capability profile of different methods under a 150M-token budget (LLama). ReAD consistently outperforms all six baselines.

Method	General	Steerability	Reasoning	Math	Code	Tool Use	LCU	Multilingual
Teacher only	48.13 $\pm$ 0.35	89.98 $\pm$ 0.40	10.51 $\pm$ 0.12	48.34 $\pm$ 0.55	88.40 $\pm$ 0.70	31.90 $\pm$ 0.45	36.20 $\pm$ 0.40	91.10 $\pm$ 0.55
Student only	31.09 $\pm$ 0.40	49.22 $\pm$ 0.55	8.72 $\pm$ 0.15	15.56 $\pm$ 0.60	72.60 $\pm$ 0.95	25.83 $\pm$ 0.60	30.40 $\pm$ 0.45	68.90 $\pm$ 0.80
Resp-SFT	38.41 $\pm$ 0.58	68.73 $\pm$ 0.94	9.55 $\pm$ 0.14	28.72 $\pm$ 1.08	79.18 $\pm$ 1.26	28.53 $\pm$ 0.82	32.69 $\pm$ 0.73	78.23 $\pm$ 1.12
CoT-SFT	38.02 $\pm$ 0.61	67.92 $\pm$ 0.89	10.14 $\pm$ 0.12	31.58 $\pm$ 0.97	78.86 $\pm$ 1.33	28.41 $\pm$ 0.79	32.83 $\pm$ 0.68	77.77 $\pm$ 1.06
Logit-SFT	39.03 $\pm$ 0.56	69.38 $\pm$ 0.86	9.70 $\pm$ 0.13	29.61 $\pm$ 1.02	79.69 $\pm$ 1.21	28.71 $\pm$ 0.77	33.01 $\pm$ 0.70	78.63 $\pm$ 1.09
Resp-KD	39.62 $\pm$ 0.53	70.98 $\pm$ 0.84	9.82 $\pm$ 0.12	30.21 $\pm$ 0.98	80.09 $\pm$ 1.18	29.09 $\pm$ 0.74	33.42 $\pm$ 0.66	79.18 $\pm$ 1.03
CoT-KD	39.12 $\pm$ 0.55	69.84 $\pm$ 0.88	10.26 $\pm$ 0.11	33.02 $\pm$ 0.92	79.81 $\pm$ 1.24	29.01 $\pm$ 0.71	33.19 $\pm$ 0.69	78.71 $\pm$ 1.01
Logit-KD	40.18 $\pm$ 0.51	72.63 $\pm$ 0.79	9.95 $\pm$ 0.12	31.36 $\pm$ 0.95	80.57 $\pm$ 1.16	29.48 $\pm$ 0.69	33.73 $\pm$ 0.64	79.97 $\pm$ 0.97
ReAD	41.80 $\pm$ 0.45	78.20 $\pm$ 0.70	10.40 $\pm$ 0.10	35.20 $\pm$ 0.80	83.90 $\pm$ 1.00	31.30 $\pm$ 0.55	35.60 $\pm$ 0.55	84.10 $\pm$ 0.90

Table 5: Capability profile of different methods under a 20M-token budget (Qwen). ReAD consistently outperforms all six **one-hot** baselines **on every capability**.

Method	General	Steerability	Reasoning	Math	Code	Tool Use	LCU	Multilingual
Teacher only	55.27 $\pm$ 0.46	92.36 $\pm$ 0.52	14.28 $\pm$ 0.18	55.74 $\pm$ 0.78	90.12 $\pm$ 0.88	36.08 $\pm$ 0.63	40.23 $\pm$ 0.57	92.34 $\pm$ 0.66
Student only	40.07 $\pm$ 0.61	70.18 $\pm$ 0.74	11.47 $\pm$ 0.22	28.26 $\pm$ 0.95	81.96 $\pm$ 1.10	30.06 $\pm$ 0.82	34.17 $\pm$ 0.71	83.14 $\pm$ 0.93
Resp-SFT	42.18 $\pm$ 0.86	72.92 $\pm$ 0.98	11.62 $\pm$ 0.24	31.42 $\pm$ 1.35	83.88 $\pm$ 1.42	30.96 $\pm$ 0.97	35.04 $\pm$ 0.88	84.36 $\pm$ 1.21
CoT-SFT	41.93 $\pm$ 0.90	72.41 $\pm$ 1.02	12.06 $\pm$ 0.21	33.12 $\pm$ 1.28	83.54 $\pm$ 1.50	30.74 $\pm$ 0.95	34.92 $\pm$ 0.91	84.09 $\pm$ 1.18
Logit-SFT	42.57 $\pm$ 0.82	73.56 $\pm$ 0.95	11.81 $\pm$ 0.23	32.24 $\pm$ 1.22	84.06 $\pm$ 1.37	31.13 $\pm$ 0.93	35.21 $\pm$ 0.86	84.58 $\pm$ 1.16
Resp-KD	42.79 $\pm$ 0.80	74.03 $\pm$ 0.92	11.93 $\pm$ 0.22	32.68 $\pm$ 1.18	84.22 $\pm$ 1.30	31.42 $\pm$ 0.90	35.44 $\pm$ 0.83	84.92 $\pm$ 1.10
CoT-KD	42.44 $\pm$ 0.83	73.37 $\pm$ 0.94	12.31 $\pm$ 0.19	34.08 $\pm$ 1.12	83.91 $\pm$ 1.28	31.02 $\pm$ 0.88	35.29 $\pm$ 0.85	84.76 $\pm$ 1.08
Logit-KD	43.06 $\pm$ 0.78	74.74 $\pm$ 0.88	12.04 $\pm$ 0.21	33.07 $\pm$ 1.10	84.55 $\pm$ 1.22	31.77 $\pm$ 0.86	35.73 $\pm$ 0.80	85.27 $\pm$ 1.02
ReAD	44.63 $\pm$ 0.73	78.57 $\pm$ 0.86	12.86 $\pm$ 0.17	36.24 $\pm$ 1.05	86.38 $\pm$ 1.18	33.41 $\pm$ 0.79	36.98 $\pm$ 0.77	87.23 $\pm$ 0.95

D ADDITIONAL EXPERIMENTS

Additional evidence showing that ReAD improves capability distillation under a fixed budget.

We compare ReAD with six capability distillation baselines. Tables 1 (in the main text) and 4 report the Llama results at $B=20\text{M}$ and $B=150\text{M}$, and Tables 5 and 6 report the corresponding Qwen results. In both budgets and for both model families, ReAD achieves the best overall capability profile and outperforms all six baselines on every capability, with the largest improvements appearing on bottleneck capabilities such as Steerability and Math while still maintaining strong gains on Code, Tool Use, and LCU. Scaling the budget from 20M to 150M strengthens all methods, but ReAD benefits more consistently, indicating that its allocation policy converts additional tokens into broader capability improvements rather than concentrating gains on a narrow subset of skills. Across model families, Qwen starts from a stronger student and teacher on most capabilities and also reaches higher absolute performance after distillation, but the relative ordering between methods is unchanged and ReAD remains the top performer, suggesting that ReAD is robust to the choice of teacher–student backbone and that its advantage is not tied to a specific LLM family. **Figure 4 visualizes the same Llama numbers as absolute gains over the initial student, next to the best one-hot baseline of each capability, which makes the per-capability margins directly comparable.**

ReAD beats SOTA baselines. To further validate ReAD against SOTA capability distillation baselines that are specifically designed for certain capabilities, we include three representative methods that are commonly used for specializing student models, one per capability: step-by-step distillation (Hsieh et al., 2023) for reasoning, Equation-of-Thought Distillation (EoTD) (Zhu et al., 2024) for math, and CodePLAN (Sun et al., 2024) for code. For a fair comparison under our setting, we adapt each baseline to our LLaMa teacher–student pair and token budget as 150M, and we distill the student using only the corresponding capability pool. We report the resulting capability profiles on all eight benchmarks using the same evaluation protocol as the main experiments in Figure 5. The results show that these capability-targeted baselines mainly concentrate gains on their target capability and exhibit weaker transfer to other capabilities, producing an imbalanced profile that matches the interaction patterns observed in our empirical study. In contrast, ReAD achieves a consistently stronger and more balanced profile across capabilities, and it is never worse than these baselines on its target capability, while improving substantially on non-target capabilities.

Significance and budget scaling of the reported numbers. Table 7 applies unpaired two-sided Welch t -tests (three seeds per method) to the comparisons reported above, using only the printed

Table 6: Capability profile of different methods under a 150M-token budget (Qwen). ReAD consistently outperforms all six baselines.

Method	General	Steerability	Reasoning	Math	Code	Tool Use	LCU	Multilingual
Teacher only	55.27 $\pm$ 0.46	92.36 $\pm$ 0.52	14.28 $\pm$ 0.18	55.74 $\pm$ 0.78	90.12 $\pm$ 0.88	36.08 $\pm$ 0.63	40.23 $\pm$ 0.57	92.34 $\pm$ 0.66
Student only	40.07 $\pm$ 0.61	70.18 $\pm$ 0.74	11.47 $\pm$ 0.22	28.26 $\pm$ 0.95	81.96 $\pm$ 1.10	30.06 $\pm$ 0.82	34.17 $\pm$ 0.71	83.14 $\pm$ 0.93
Resp-SFT	48.02 $\pm$ 0.92	82.07 $\pm$ 1.04	13.18 $\pm$ 0.20	41.26 $\pm$ 1.38	86.92 $\pm$ 1.40	33.29 $\pm$ 0.96	37.42 $\pm$ 0.88	88.07 $\pm$ 1.20
CoT-SFT	48.29 $\pm$ 0.88	82.61 $\pm$ 0.98	13.72 $\pm$ 0.18	44.43 $\pm$ 1.30	86.58 $\pm$ 1.45	33.05 $\pm$ 0.94	37.28 $\pm$ 0.90	87.96 $\pm$ 1.18
Logit-SFT	48.46 $\pm$ 0.86	83.14 $\pm$ 0.95	13.39 $\pm$ 0.19	42.68 $\pm$ 1.26	87.05 $\pm$ 1.32	33.47 $\pm$ 0.92	37.63 $\pm$ 0.86	88.31 $\pm$ 1.12
Resp-KD	48.78 $\pm$ 0.84	83.76 $\pm$ 0.90	13.51 $\pm$ 0.18	43.62 $\pm$ 1.22	87.37 $\pm$ 1.26	33.92 $\pm$ 0.88	38.02 $\pm$ 0.82	88.64 $\pm$ 1.06
CoT-KD	49.05 $\pm$ 0.82	84.28 $\pm$ 0.88	13.87 $\pm$ 0.17	45.17 $\pm$ 1.18	87.11 $\pm$ 1.20	33.66 $\pm$ 0.86	37.91 $\pm$ 0.84	88.53 $\pm$ 1.02
Logit-KD	49.37 $\pm$ 0.80	85.19 $\pm$ 0.84	13.62 $\pm$ 0.18	44.07 $\pm$ 1.16	87.63 $\pm$ 1.18	34.26 $\pm$ 0.84	38.33 $\pm$ 0.78	88.92 $\pm$ 0.98
ReAD	51.53 $\pm$ 0.76	88.04 $\pm$ 0.90	14.03 $\pm$ 0.15	47.06 $\pm$ 1.05	89.04 $\pm$ 1.12	35.52 $\pm$ 0.78	39.47 $\pm$ 0.74	90.54 $\pm$ 0.92




Figure 4: Points gained over the initial student on each capability at 20M (left) and 150M (right) for ReAD and for the best one-hot baseline of that capability, read off Tables 1 and 4. Error bars are the standard deviation over three seeds of the distilled student’s score. Values are absolute percentage points, so small changes stay small.

means and standard deviations, with Holm correction within each block. At 150M every gain over the allocation baselines of Table 2(b) remains significant after correction, and so do the Steerability and Multilingual gains over the best one-hot baseline; at 20M, with three seeds, most per-capability differences are individually significant only before correction. The XSTest safe-refusal improvements are not significant at this sample size. Table 8 reports how much each method gains per additional 10M tokens between the two budgets: ReAD has the largest marginal gain on six of the eight capabilities and on average, which is the budget-efficiency claim of RQ2 stated in absolute points.

Local proxy validation. The proxy validation supports ReAD’s use of local action ranking. Predicted next-step gains align with observed gains, especially for small budget moves, which matches ReAD’s stepwise allocation regime. The requirement vectors are stable across runs, and the reward regressors are accurate enough to guide the uncertainty-aware bandit updates. (Moved from the main text.)

Held-out Downstream Transfer. To test whether ReAD transfers beyond capability-aligned benchmark tasks, we evaluate on XSTest using the same requirement identifier and allocation pipeline without adding a new capability head. Lower safe-refusal and higher unsafe-refusal scores are better.

E RELATED WORK

Capability Distillation for Large Language Models. Recent work on capability distillation has focused on transferring specific capabilities such as instruction following and context manipulation (Taori et al., 2023; Chiang et al., 2023; Peng et al., 2023; Xu et al., 2023a), thinking pat-




Figure 5: ReAD beats SOTA baselines that are specifically designed for reasoning, code, and math capabilities consistently, at the 150M-token budget.

Table 7: Statistical significance of the reported comparisons, computed from the means and standard deviations in Tables 1, 4, 2(b), and 9 (three seeds; unpaired two-sided Welch t -test on summary statistics). Δ is ReAD minus the competitor in that row; p values are Holm-corrected within each block. *: $p < 0.05$, **: $p < 0.01$. A paired test on per-seed scores would be more powerful; per-seed scores will be released with the code.

Comparison	Competitor	20M		Competitor	150M	
		Δ	t		Δ	t
General	Logit-KD	+1.00	2.9	Logit-KD	+1.62	4.1
Steerability	Logit-KD	+3.82	7.4*	Logit-KD	+5.57	9.1**
Reasoning	CoT-KD	+0.39	3.8	CoT-KD	+0.14	1.6
Math	CoT-KD	+1.72	3.0	CoT-KD	+2.18	3.1
Code	Logit-KD	+2.44	2.9	Logit-KD	+3.33	3.8
Tool Use	Logit-KD	+2.41	4.4	Logit-KD	+1.82	3.6
LCU	Logit-KD	+1.86	4.6	Logit-KD	+1.87	3.8
Multilingual	Logit-KD	+2.23	3.1	Logit-KD	+4.13	5.4*
Avg. utility (Table 2b)	Greedy one-step	+1.01	3.1	Greedy one-step	+1.20	5.1**
Avg. utility (Table 2b)	Grid-searched	+1.16	3.5	Grid-searched	+1.40	5.8**
Avg. utility (Table 2b)	Task-static mix	+1.33	4.0	Task-static mix	+1.69	6.8**
Avg. utility (Table 2b)	Uniform static	+2.10	5.9*	Uniform static	+2.48	9.4**
XSTest safe-refusal ↓	best baseline	-1.3	-2.1	best baseline	-1.3	-2.5
XSTest unsafe-refusal ↑	best baseline	+2.6	3.1*	best baseline	+2.6	3.7*

terns (Zhang et al., 2024; Cheng et al., 2023), or text and code generation capability (Zhang et al., 2025; Li et al., 2023). These methods typically steer a teacher LLM via prompts or rationales, generate distillation data or logits, and fine-tune the student on that evidence. However, the majority of methods operate under the assumption that distilling one target capability or domain is sufficient and rarely analyze how that focused training influences broader capability interactions within the student model. To the best of our knowledge, ReAD is the first framework that explicitly accounts for capability interdependence during the distillation.

Distillation objectives, data mixtures, and forgetting. Beyond capability-specific recipes, LLM distillation methods differ in where the training sequences come from. Sequence-level KD and its descendants train on teacher-generated outputs (Hinton et al., 2015; Kim & Rush, 2016), whereas on-policy methods train on student-generated sequences scored by the teacher, which reduces exposure bias (Gu et al., 2024; Agarwal et al., 2024; Ko et al., 2024). ReAD is off-policy in this sense and treats the objective as a plug-in: its contribution is the allocation of a fixed budget across capabilities, not a new loss. A separate line of work optimizes the *mixture* of training data: DoReMi (Xie et al., 2023) and RegMix (Liu et al., 2024) learn static domain weights before, or without observing, the target model’s training trajectory, whereas ReAD re-estimates the mixture online from the student’s

Table 8: Marginal gain from the 20M to the 150M budget, in points per additional 10M tokens, computed from Tables 1 and 4 as $(s_{150M} - s_{20M})/13$. Bold: largest per column.

Method	General	Steer.	Reason.	Math	Code	Tool	LCU	Multi.	Mean
Resp-SFT	0.39	1.22	0.04	0.89	0.39	0.12	0.11	0.63	0.48
CoT-SFT	0.38	1.21	0.05	1.00	0.39	0.12	0.13	0.61	0.49
Logit-SFT	0.42	1.26	0.04	0.92	0.40	0.13	0.12	0.65	0.49
Resp-KD	0.45	1.34	0.05	0.95	0.45	0.15	0.14	0.67	0.52
CoT-KD	0.44	1.30	0.05	1.05	0.44	0.15	0.13	0.66	0.53
Logit-KD	0.47	1.42	0.05	1.01	0.46	0.16	0.15	0.70	0.55
ReAD	0.52	1.55	0.03	1.09	0.53	0.11	0.16	0.85	0.60

Table 9: Held-out XSTest evaluation. “Best baseline” denotes the best of the six SFT/KD baselines for each metric and budget separately.

Method	20M safe ↓	20M unsafe ↑	150M safe ↓	150M unsafe ↑
Best baseline	15.1 ± 0.77	75.5 ± 1.09	13.1 ± 0.69	79.0 ± 0.88
ReAD	13.8 ± 0.72	78.1 ± 0.97	11.8 ± 0.58	81.6 ± 0.83

measured capability profile and the remaining budget. Finally, the off-target degradation we measure is related to catastrophic forgetting in continual fine-tuning (Kirkpatrick et al., 2017; Luo et al., 2023); ReAD adds no explicit anti-forgetting regularizer but penalizes measured spillover on task-essential capabilities directly in its reward (Eq. 9).

F BROADER IMPACT

ReAD aims to make capability distillation more budget-efficient by allocating limited training tokens across interdependent LLM capabilities. A positive impact is that smaller student models can become more useful under fixed compute budgets, potentially reducing deployment cost and improving access to LLM-based systems in resource-constrained settings. ReAD also explicitly monitors cross-capability transfer, which can help practitioners diagnose when improving one capability degrades other capabilities that a deployment depends on.

At the same time, more efficient distillation may lower the cost of producing specialized models with dual-use capabilities, such as code generation, tool use, or instruction following. Distilled models may also inherit biases, unsafe behaviors, privacy risks, or other limitations from teacher models and training data. In addition, if the capability probes or task requirement vectors are misaligned with the intended deployment objective, ReAD may optimize benchmark utility while overlooking safety-, fairness-, or robustness-relevant behaviors.

We therefore view ReAD as a research framework rather than a deployment-ready safety guarantee. For sensitive applications, practitioners should combine ReAD with task-specific safety evaluation, robustness and bias audits, data filtering, red-team testing, and monitoring of off-target capability changes. Any release or deployment of distilled models should follow the usage restrictions of the underlying models and include appropriate documentation of intended use, limitations, and the safety evaluations that were actually performed.

G LIMITATION AND FUTURE WORK

While ReAD improves budget efficiency by using low-cost capability monitoring and adaptive allocation, it still relies on two practical components that can limit performance. First, the task requirement estimator and the probe suite may be imperfect: if the inferred essential capabilities or probe signals are misaligned with true task utility, ReAD can allocate budget suboptimally. Second, our capability set and data generator are necessarily finite and template-driven, which may not cover all behaviors a real task depends on, especially for highly domain-specific or long-horizon tasks. Future work can address these limitations by learning more robust and transferable task requirement models, designing probe suites with stronger coverage and better calibration to task utility, and expanding capability definitions and data generation beyond templates (e.g., automatic discovery of capability axes and open-ended curriculum generation). Another promising direction is to jointly

optimize allocation and student training dynamics (e.g., optimizer schedules and parameter-efficient adapters) under the same budget, and to extend ReAD to settings with distribution shift, multi-task deployment, or continual updates where requirements and transfer patterns evolve over time.